

Intel Arch SIG Meeting

Scalable IOV Introduction

Jane Lv & Jun Tian (SATG/SSE Linux OS Enabling team)

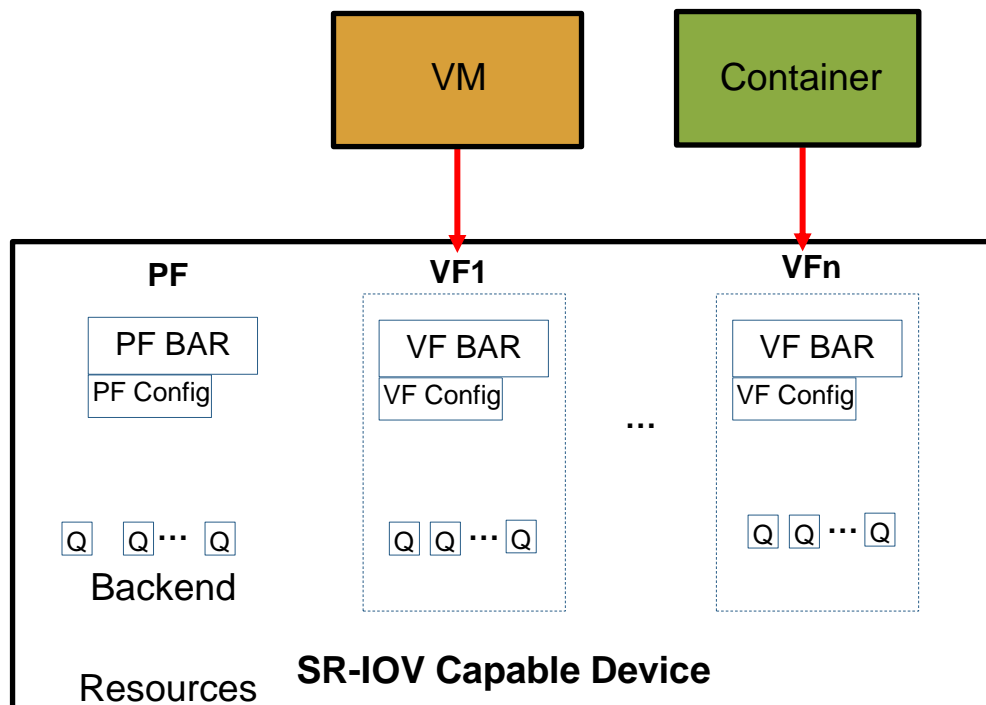


Agenda

- SIOV Introduction
- Upstream Refactoring

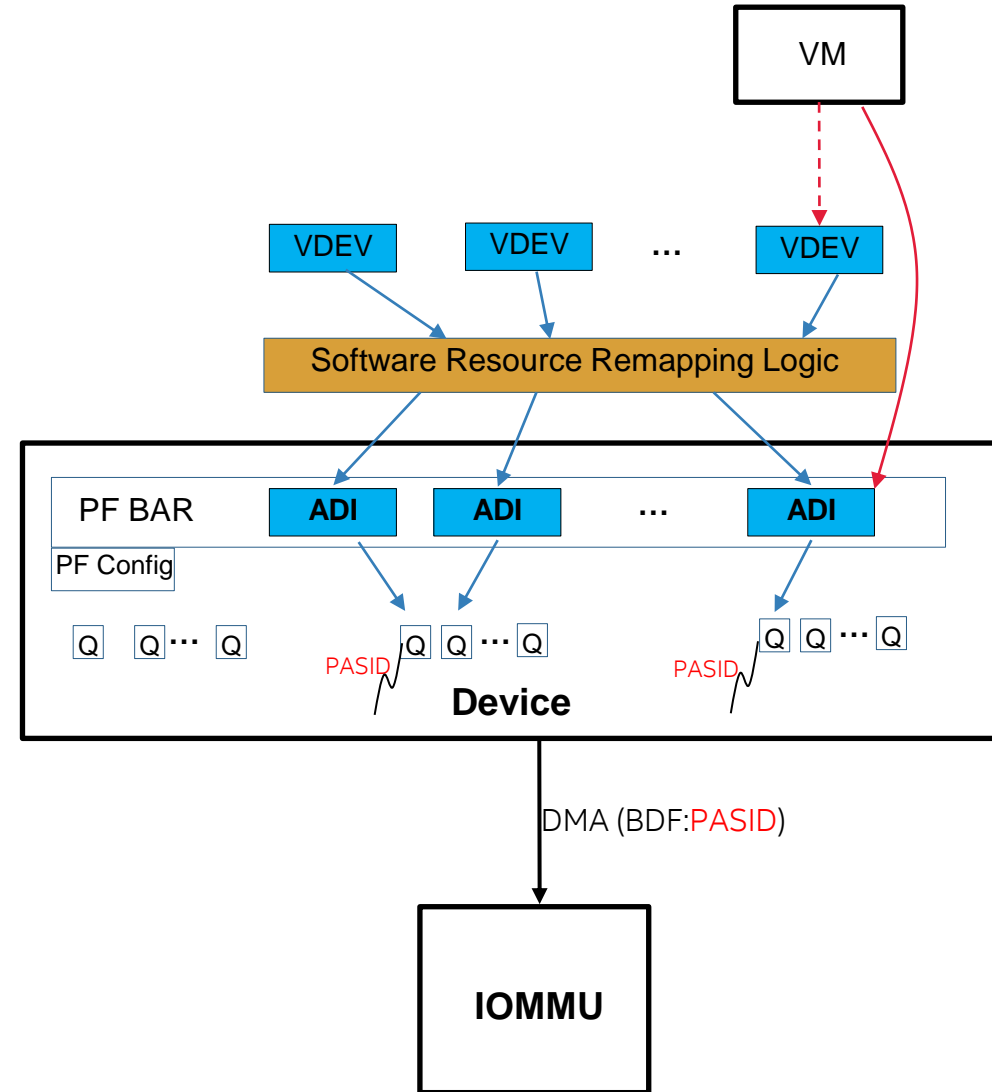
SIOV Introduction

SIOV vs SR-IOV



PCI Single Root I/O Virtualization (SR-IOV)

- Physical Function & Virtual Function
- VF directly assignable to VM/container



Scalable IOV

- Assignable Device Interface (ADI), Queues, Q pairs, contexts with PASID
- PASID-granule DMA isolation
- Compose ADIs into Virtual Device (Vdev)

SIOV Advantages

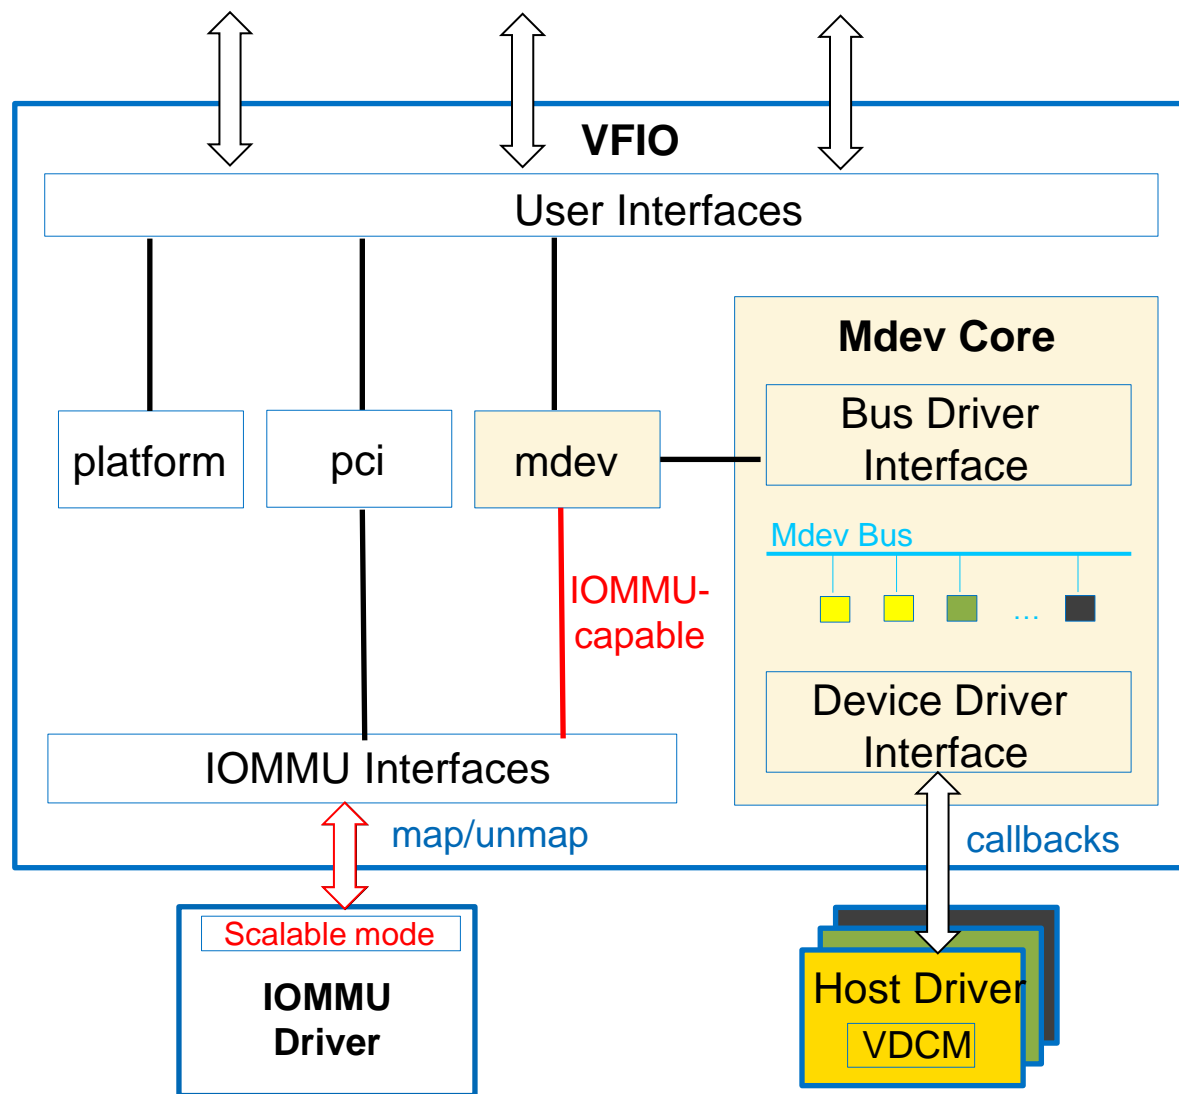
Usage Point-of-view

- **Hyper-scale usages**
 - E.g. scale to 1000+ VMs/containers
- **Dynamic resource management**
 - E.g. user-defined sharing granularity, over-provisioning, etc.
- **Composability**
 - E.g. VM live migration, snapshot, etc.

Technical Point-of-view

- **A hardware-assisted mediated pass-through architecture**
 - Slow-path operations emulated by SW
 - Fast-path resources dynamically provisioned for direct access
 - Hardware-enforced DMA isolation between fast-path resources
- **Finer-grained device sharing than SR-IOV**
 - Think about each TX/RX queue pair is now assignable
- **Supports any type of devices**
 - Integrated or discrete
- **Utilizes existing PCIe capabilities**
 - e.g. Process Address Space ID (PASID)

SIOV Framework (OOT)



■ IOMMU-capable mdev

- ✓ Link to `iommu_domain` (tagged by PASID)
- ✓ Allow PASID-granule iommu map/unmap

■ Opt-in by VDCM

- ✓ When a mdev is created

SIOV Concepts

- Work Queues (WQ) - On device storage to queue descriptors to the device. Requests are added to a WQ by using new instructions to write to the memory mapped “portal” associated with each WQ.
 - Dedicated WQ (DWQ) - A single client owns this exclusively and can submit work to it.
 - A client using DWQ submits work descriptors using the **MOVDIR64B** instruction. Posted write, do not exceed the configured length of WQ
 - Shared WQ (SWQ) - Multiple clients can submit work to the SWQ.
 - Clients using shared work queues submit work descriptors using either ENQCMDS (from supervisor mode) or **ENQCMD** (from user mode).
 - These instructions indicate via the EFLAGS.ZF bit whether the request was accepted.
 - A SWQ requires PASID and can only run with SVA support

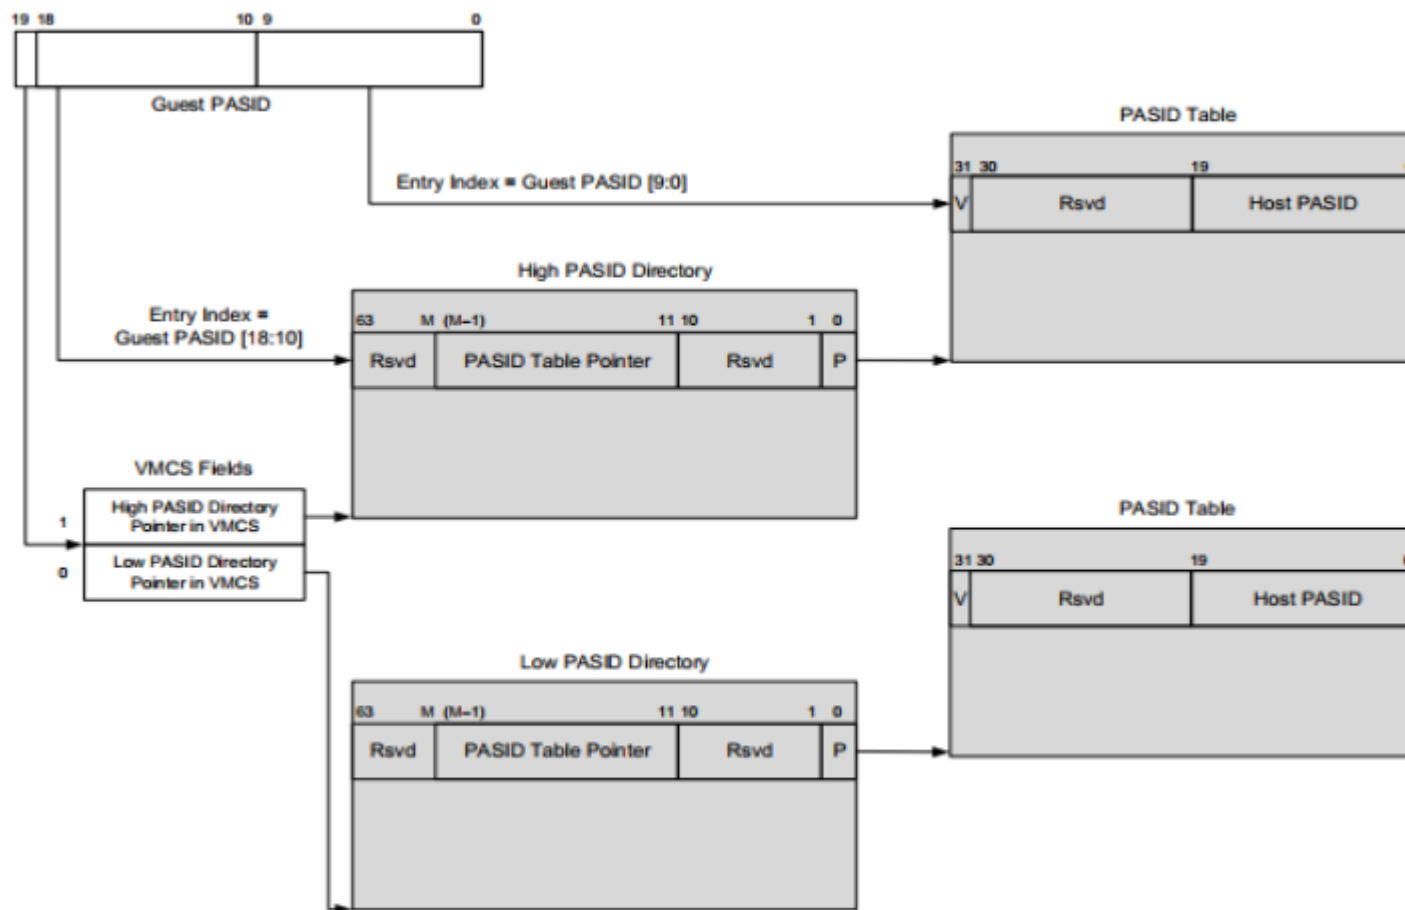
Shared WQ in VM ← PASID in VM ← PASID Virtualization ← vIOMMU

ENQCMD/ENQCMLS

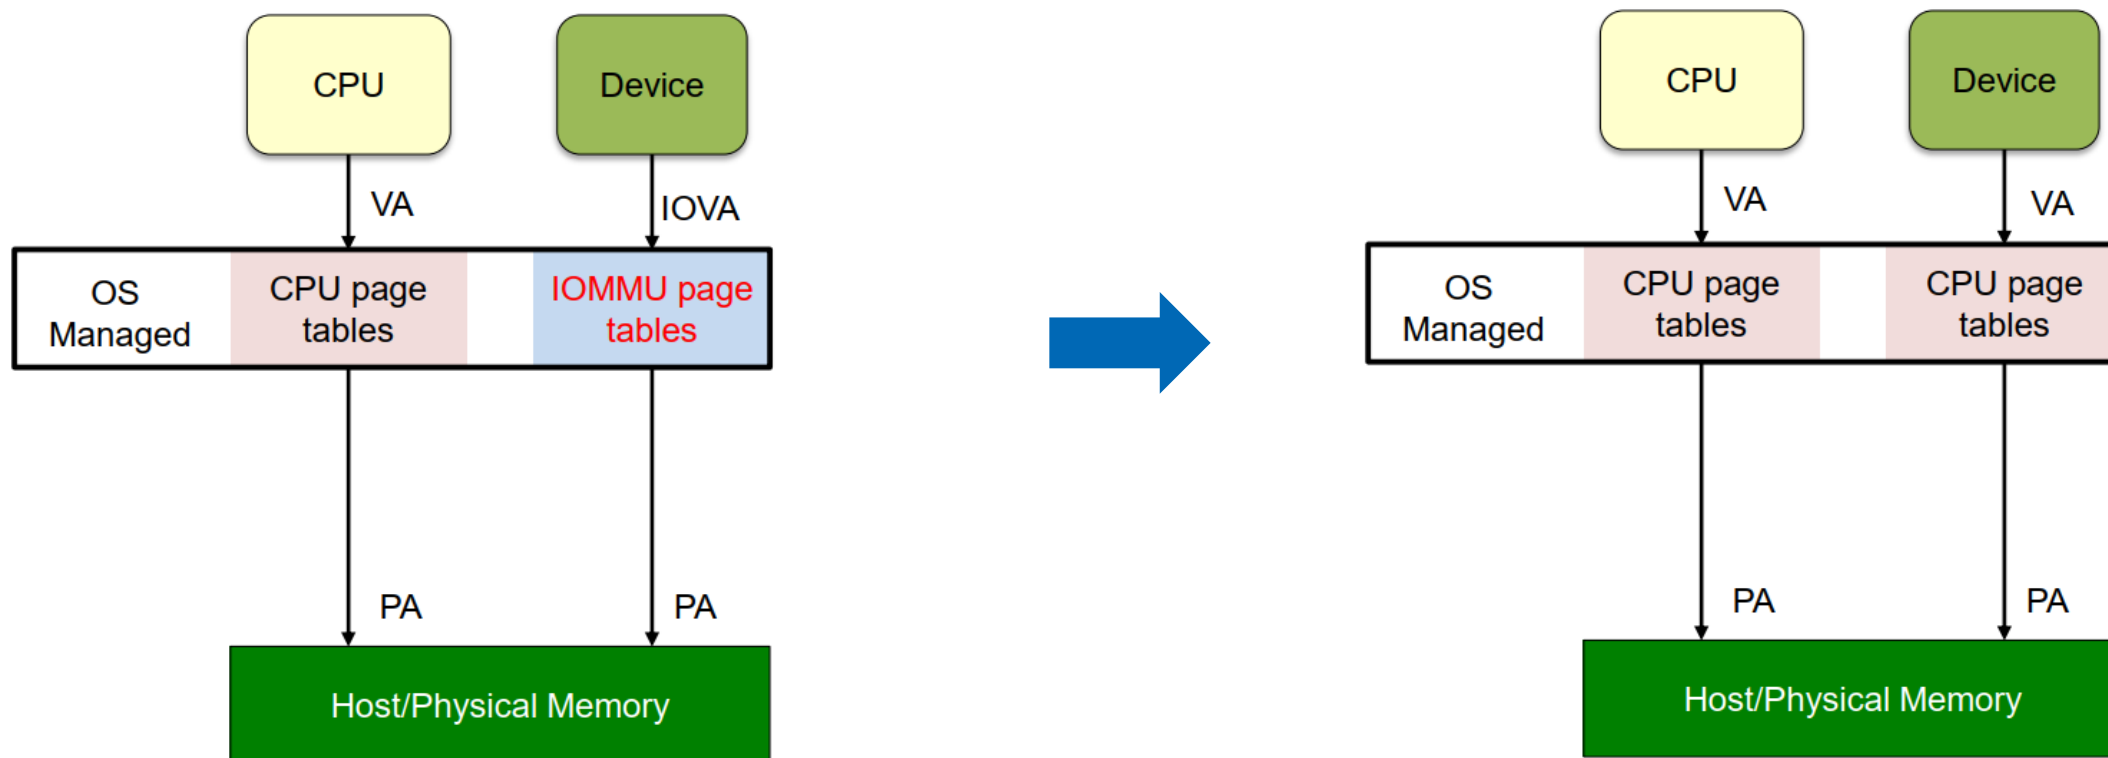
- New instruction on Intel® Platforms
- Atomically submit a work with PASID
 - Obtains PASID from IA32_PASID MSR
 - Enqueue store 64B command to enqueue register in device MMIO
- IA32_PASID is managed by XSAVE as PASID supervisor state
- Non-Posted instruction which carries back a status
 - ZF flag indicates if the command was accepted by device
 - Allows user to retry
- ENQCMLS for Supervisor
 - PASID from command data

ENQCMD in VM

- New feature in VMX on Intel® Platforms
- Use PASID Translation Table for guest PASID to host PASID translation
- Trigger VM-Exit if fails to translate guest PASID



Shared Virtual Address




Provide system and device memory unified programming model

SVA Concepts

- Process Address Space ID (PASID)
 - Identify process address space
- First-level translation
 - DMA requests with PASID
 - For SVA transaction from endpoint device (GIOVA/GVA->GPA)
- Second-level translation
 - DMA requests without PASID
 - For normal DMA transaction from endpoint device (GPA->HPA)

Translation Types

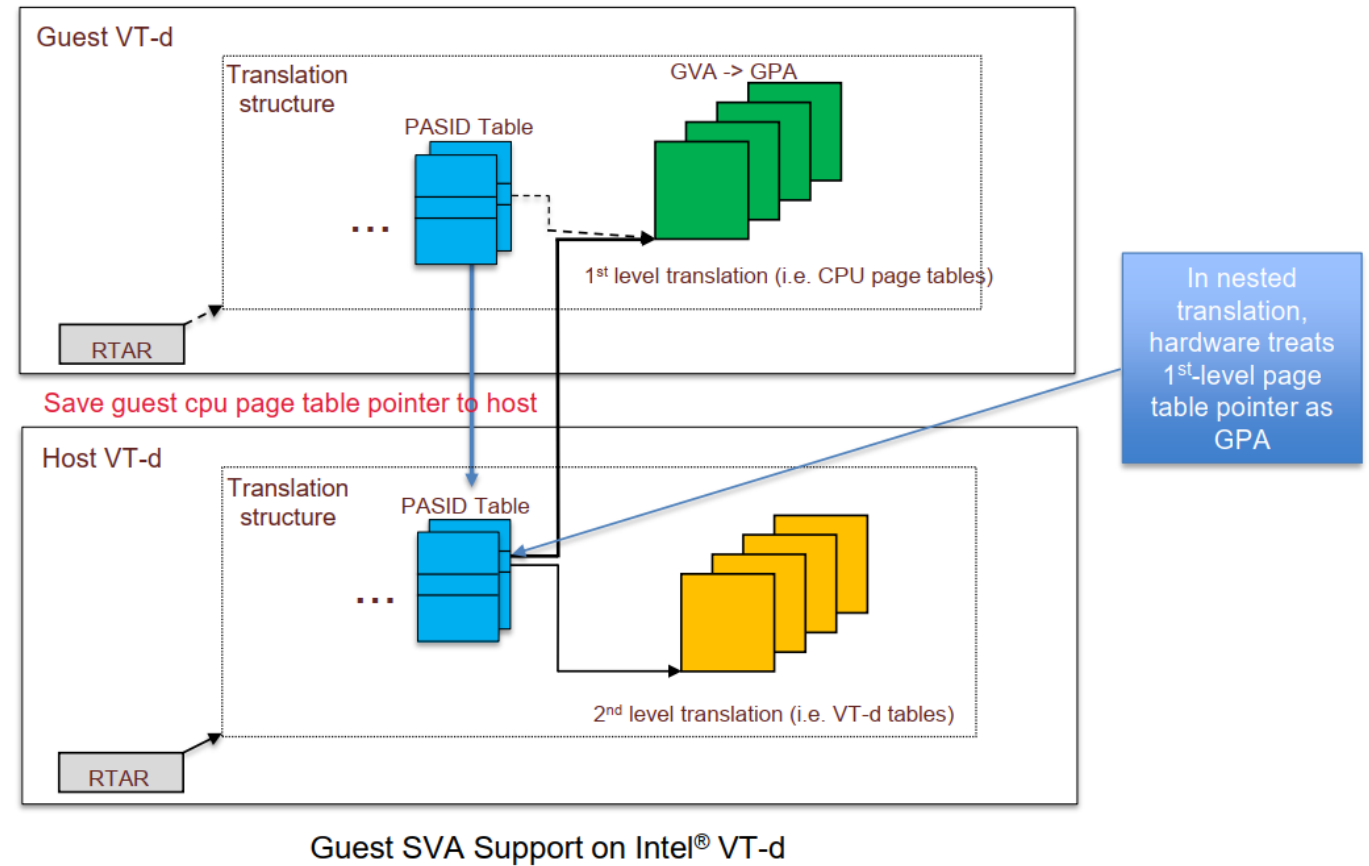
- First-Level translation
- Second-Level translation
- Nested translation - Virtualization environment
 - Need a vIOMMU with SVA capability)
 - Vendor specific
- Pass-Through (address translation bypassed)



SVA in VM

Enable SVA in VM a.k.a vSVA

- vIOMMU emulation in QEMU
- VFIO for programming host IOMMU
- IOMMU Driver with new API to expose VFIO for guest SVA (vSVA)
- Enable nested translation
- PASID virtualization



Native SVA

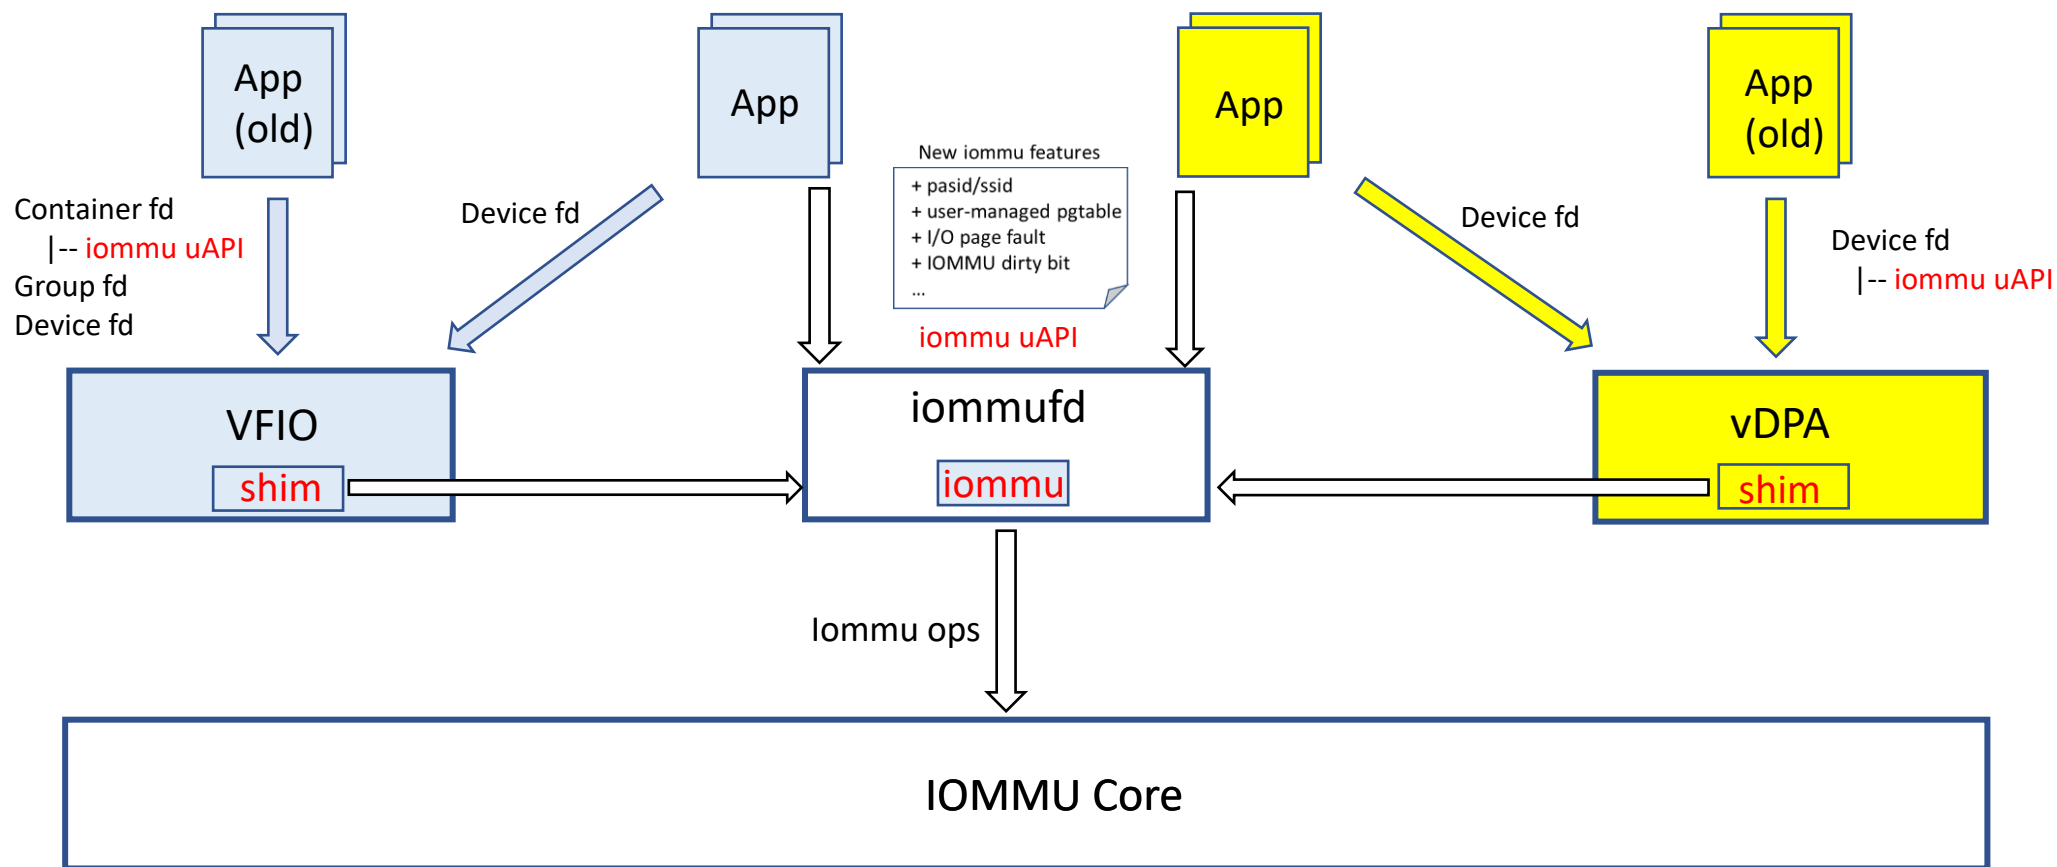
- Intel iommu driver has SVA support since kernel 5.0
- Generic SVA framework is developed in community
 - Generic kernel API is introduced (iommu_sva_bind/unbind_device())
 - IOASID is introduced to manage PASID (merged in 5.5)
 - Uacce is introduced to support userspace usage (merged in 5.7)

Intel VT-d Enhancement

- Scalable mode DMA remapping
 - PASID granule 1st-level, 2nd-level, nested and pass-thru
 - PASID table now two-level structure
 - Cover both Scalable IOV and SVA usages
 - Extended Context (ECS) is deprecated
- Access/Dirty (A/D) bits in 2nd-level
 - Assist dirty memory tracking in live migration

Upstream Refactoring

Refactoring for a Unified Framework (iommuufd)



<https://github.com/jgunthorpe/linux/commits/iommuufd>

